

Department of Biostatistics
Erasmus University Medical Center
PO Box 2040, 3000 CA Rotterdam
the Netherlands

February 9, 2018

Professor Michael J. Daniels
Department of Statistics
University of Florida
Gainesville, FL, 32611-8545
USA

Dear Professor Daniels,

We are writing to you with respect to the manuscript #BIOM2017609M, titled “Personalized schedules for surveillance of low risk prostate cancer patients” submitted to *Biometrics* and the reports we received after its review. We would like to thank you for giving us the opportunity to submit a revised version of our paper that tackles the weaknesses of the previous version.

Following the recommendations from the Reviewers, we have made several changes in the revised version of the manuscript. In particular, we have invested more in the new version in To this end The previous version was 25 pages long and you have asked us to reduce that to ... pages (including title, authorship, abstract, body of manuscript, acknowledgments, references). While according to the suggestions of the Reviewers we have included new pieces of information, we have managed to reduce the length of the paper to 21.5 pages. We hope that this is acceptable.

Please find enclosed a detailed point-by-point response to the Reviewers’ comments.

Yours sincerely,

the Authors

Response to AE's Comments

We would like to thank the Associate Editor for his/her constructive comments, which have allowed us to considerably improve our paper. The main differences of the new version of the manuscript compared to the previous one can be found in Sections 2, 3 and 5. In addition, changes regarding the specific comments have been made throughout the text.

With respect to the AE's note about how delaying active treatment relates to survival outcome, the biopsies are already conducted for the patients according to the PRIAS schedule. Thus we cannot test the effect of delay in detection Gleason reclassification (GR) due to personalized schedules on the prostate cancer (PCa) specific mortality, unless they are applied on the patients. It may however be possible to compare the effect of delay in detection of GR due to PRIAS (less frequent biopsies) and annual schedules (more frequent biopsies), on the PCa specific mortality. However this would be out of the scope of our current focus on personalized schedules. In addition, multiple studies have reported small PCa specific mortality in low risk patients enrolled in active surveillance programs (Klotz et al., 2009; Loeb et al., 2016; Tosoian et al., 2011). That is, less frequent schedules may be useful in such scenarios. For example, for slowly progressing patients (subgroup G_3) in our simulation study, we observed that even a personalized schedule which conducts on average two biopsies leads to an average delay of 10 months in detecting GR. This is only four months more delay than that of the annual schedule. Given, the low PCa mortality a relative difference of four months may not be that bad an alternative.

Response to 1st Referee’s Comments

We would like to thank the Referee for his/her constructive comments, which have allowed us to considerably improve our paper. The main differences of the new version of the manuscript compared to the previous one can be found in Sections 2, 3 and 5. In addition, changes regarding the specific comments have been made throughout the text. You may find below our response to the specific issues raised.

1. With regards assumption of normality of random effects, joint models (JM) have been shown to be quite robust to random effects misspecification. More specifically, Huang, Stefanski, and Davidian (2009) and Rizopoulos, Verbeke, and Molenberghs (2008) have shown that unless the number of repeated measurements per person are extremely small, such misspecification only and trivially affects the standard errors. On the other hand in our dataset we have a mean of 8.7 measurements per person.

With regards to the assumption of normality of error term, we conducted residual diagnostics to check this assumption. The quantile-quantile (q-q) plot of subject specific residuals in Figure 1 shows that a long tailed distribution for errors is more plausible than the normal distribution. Based on this result, we further fitted two JMs with t-distributed errors, with 4 and 3 degrees of freedom (df), respectively. We found that the model with t-distributed (df=3) errors satisfied the distributional assumptions the best (see Figure 1).

We then compared the model with assumption that errors are normally distributed and the model with assumption that errors are t-distributed. To this end, the fitted marginal \log_2 PSA profile for a hypothetical patient with age 70 years using the two models is shown in Figure 2. We also compared the subject specific fitted \log_2 PSA profiles for 9 randomly selected patients (each with more than 3 observations). A seed of 2017 (year of submission of the article) was used to sample these patients from the PRIAS dataset sorted by patient ID. Lastly, for the two models, Table 1 shows the association parameters. We can see that the association between the hazard of GR and slope of \log_2 PSA is stronger in the model with t-distributed (df=3) errors. In light of the better fit of the model with t-distributed (df=3) errors, we have updated the parameters estimates in Web Appendix C of the updated supplementary material.

Since the slope association between \log_2 PSA levels and hazard of Gleason reclassification (GR) in the model with t-distributed (df=3) has become stronger, we expect our schedules to become slightly more sensitive towards increase in \log_2 PSA velocity. This however may




Figure 1: Quantile-quantile plots of subject specific residuals obtained from joint models with assumption of normally distributed errors, and t-distributed (df=3) errors, fitted to the PRIAS data set.

Table 1: Relative risk sub-model estimates for association parameters between hazard of GR and slope of \log_2 PSA levels. Mean and 95% credible interval (CI) are presented for fits obtained from joint model with assumption of normal distributed errors, and t-distributed (df=3) errors.

Error distribution	\log_2 PSA association [95% CI]	Slope(\log_2 PSA) association [95% CI]
t-distribution (df=3)	-0.004 [-0.119, 0.117]	2.888 [2.318, 3.452]
Normal distribution	-0.049 [-0.172, 0.078]	2.407 [1.791, 3.069]

also depend on the type of personalized schedule. For example, we compared the personalized schedule based on dynamic risk of GR using the two different models for the three demonstration patients, and observed trivial differences. This is however due to the fact that average risk (averaged over all time points) taken by dynamic risk of GR is not very high (5.3%). However quantiles corresponding to 50% risk (median time of GR) may differ by a bigger margin depending upon the profile of the patient (same for expected failure time). For example, we can see in Figure 4 and Figure ?? that the third demonstration patient has a consistent profile, with quite slow rise in PSA. Consequently the effect of the increased \log_2 PSA




Figure 2: Fitted marginal 10 year \log_2 PSA profile with 95% credible interval (CI), for a hypothetical patient who was included in AS at the age of 70 years. Fits were obtained from joint models with assumption of normal distributed errors, and t-distributed (df=3) errors. The darker shaded region indicates the overlap in the two CI intervals, as well as demarcates the two sets of CIs.

slope association parameter does not affect the schedule much for this patient. Similar results are observed for the first demonstration patient when the PSA consistently remains low over nearly three years, starting at year two. Lastly, this can also be seen for the second demonstration patient wherein the schedules differ by a large margin when PSA rises very quickly. However the gap becomes slightly smaller after a negative biopsy indicating that GR is unlikely in near future. Thus we expect that the sensitivity of the schedule based on expected failure time is within acceptable boundaries for slowly progressing (low risk) patients.




Figure 3: Fitted versus observed \log_2 PSA profiles for 9 randomly selected patients. Fits were obtained from joint models with assumption of normal distributed errors, and t-distributed (df=3) errors. The vertical line with a dot-dash pattern shows the time of the latest biopsy. The fitted profiles utilize information from both the observed PSA levels and time of latest biopsy.

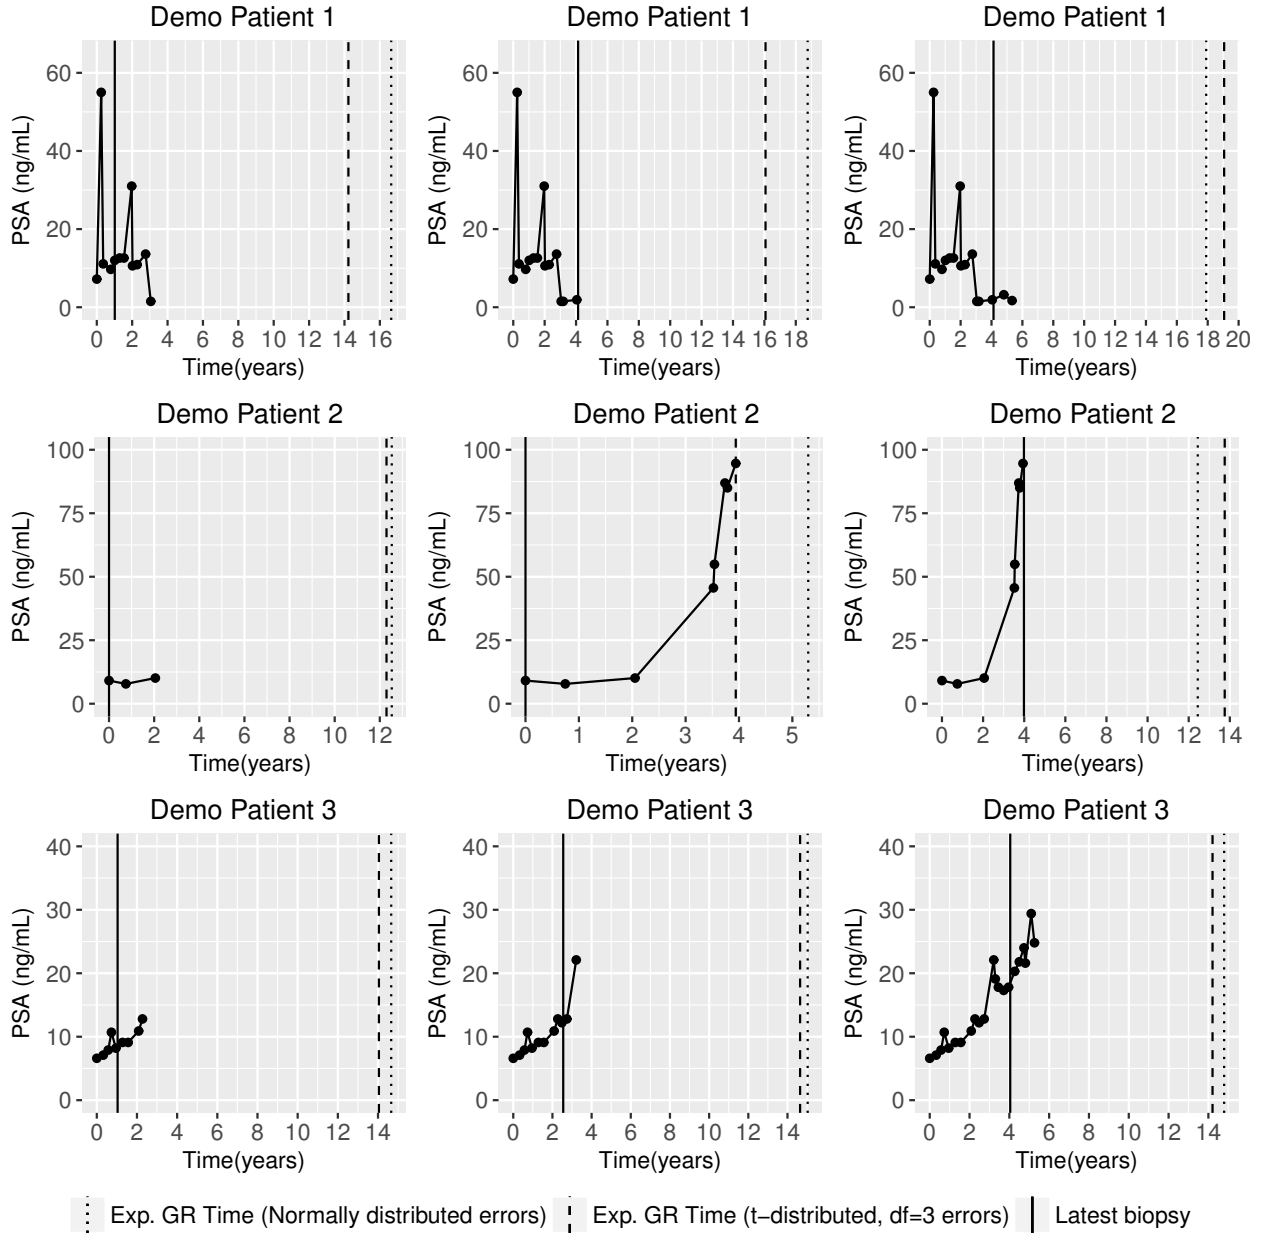


Figure 4: Dynamic expected failure time for the three demonstration patients at three different follow-up times, using joint models with assumption of normal distributed errors, and t-distributed (df=3) errors.

- The three patients were chosen on the basis of specific characteristics of their data. This is because we wanted to demonstrate three main features of the personalized schedules. For

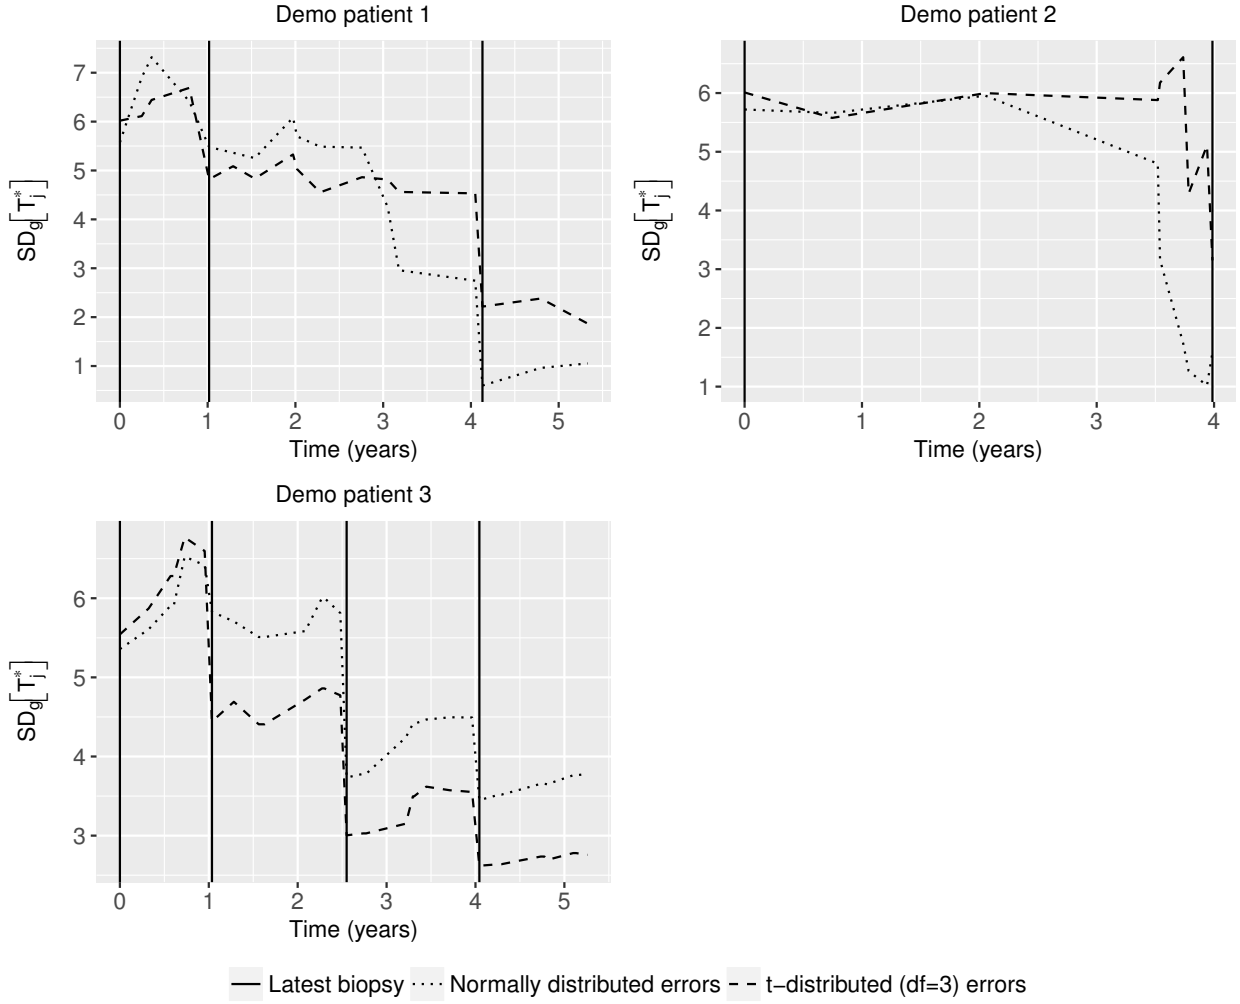


Figure 5: Dynamic variance of the posterior predictive distribution of event time for the three demonstration patients at three different follow-up times, using joint models with assumption of normal distributed errors, and t-distributed (df=3) errors.

example, the first demonstration patient had many repeat biopsies, and thus via his profile we show how the variance of posterior predictive distribution of GR time decreases with each biopsy. Via the second demonstration patient we show how the schedules change with changes in PSA alone (no repeat biopsies). Whereas, via the third demonstration patient we show how the schedules work when information from PSA and repeat biopsies is not in concordance with each other. We would like to also mention that it is not the case that we conducted an exhaustive search to purposefully select only these three patients.

With regards to conducting cross-validation on real data, and to compare the true GR time of PRIAS patients who obtained GR, with the time proposed by personalized schedules, this is not possible for the following reason. For patients in PRIAS we only know the interval $l_i < T_i^* \leq r_i$ in which GR occurred and not the true GR time T_i^* . On top of that this is known only for 707 out of 5267 patients, and the rest are right censored. That is, in either case we cannot calculate the offset $T_i^S - T_i^*$ of our schedule, where $T_i^S > T_i^*$ is the time of the last biopsy at which GR is detected. The simulation study is our attempt to obviate this problem, because there we know the true event time T_i^* for each of the patients.

3. We assume that there are three equal sized subgroups G_1 , G_2 and G_3 of patients in the population, differing in the baseline hazard of GR. This was done because we wanted to test the performance of different schedules for a population with a mixture of patients, namely those with faster progressing PCa, as well as those with slowly progressing PCa. As advised by the referee this can be modeled using a stratified modeling approach. In the current case, this corresponds to the use of latent class JMs (Proust-Lima et al., 2014). However this approach requires either knowing the number of subgroups in advance, which is unrealistic, or fitting multiple models to detect the correct number of subgroups. The latter would have been out of the scope of our paper. The alternative that we used was based on the fact that the baseline hazard in the simulated population corresponds to a mixture Weibull density (Razali and Al-Wakeel, 2013). We model the log baseline hazard flexibly using P-splines. The mean of the fitted log baseline hazard between 0.1 years and 8 years (mean third quartile in the simulated progression times is 6.15 years), and 95% confidence interval obtained from the 500 simulations is shown in Figure 6 below. It can be seen that the fit is quite close to the theoretical baseline hazard. We have added this figure and the corresponding summary in Section [section nr] of Supplementary material.
4. As noted by the referee, indeed PRIAS switches to the annual schedule if a patient's PSA doubling time (PSA-DT), measured as the inverse of the slope of the regression line through the base two logarithm of PSA values, is less than 10 years. In this regard, the joint model allows the schedule to depend upon the observed PSA values (e.g., via PSA-DT). This is because the parameters are estimated using a full likelihood approach (Tsiatis and Davidian, 2004). To show this, consider the following full general specification of the JM that we use. Let \mathbf{y}_i denote the $n_i \times 1$ vector of PSA measurements for the i -th patient, and l_i, r_i denote the two time points of the interval in which GR occurs for the i -th patient. In addition let T_i^S and \mathcal{V}_i denote the schedule of biopsies and schedule of PSA measurements, respectively.




Figure 6: Theoretical log baseline hazard of the simulated population versus mean of the fitted log baseline hazard. The 95% confidence interval for the fitted log baseline hazard is obtained from the 500 simulations.

Under the assumption that both of these schedules may depend upon only the observed \mathbf{y}_i , the joint likelihood of all four processes is given by:

$$p(\mathbf{y}_i, l_i, r_i, T_i^S, \mathcal{V}_i \mid \boldsymbol{\theta}, \boldsymbol{\psi}) = p(\mathbf{y}_i, l_i, r_i \mid \boldsymbol{\theta}) \times p(T_i^S, \mathcal{V}_i \mid \mathbf{y}_i, \boldsymbol{\psi}) \quad (1)$$

From this decomposition we can see that even if the processes T_i^S and \mathcal{V}_i may be determined from \mathbf{y}_i , if we are interested in the parameters $\boldsymbol{\theta}$ of the joint distribution of longitudinal and event outcome, we can maximize the likelihood based on the first term and ignore the second term. In other words, the second term will not carry information for $\boldsymbol{\theta}$.

In order to demonstrate this we simulated a dataset with 750 patients. The true event times T_i^* for these patients were generated using parameters from a joint model fitted to the PRIAS dataset. However this joint model did not include association between velocity of log PSA values and hazard of GR. That is, the hazard of GR $h_i(t)$ at any time t depends only on the underlying log PSA value $m_i(t)$ at that time. Furthermore, for these patients we used the schedule of PRIAS to generate the interval $l_i \leq T_i^* \leq r_i$ in which GR is detected. Thus the observed data for i -th patient is $\{\mathbf{y}_i, l_i, r_i\}$. Our aim is to show that if there is no association between $h_i(t)$ and velocity of log PSA value $m'_i(t)$, then even though the biopsy schedule depends on PSA-DT (which is a crude measure of PSA velocity), a JM fitted with both value and velocity associations will have an insignificant velocity association. In the fitted JM we found the value association (95% credible interval in brackets) to be 0.182 [0.090, 0.274], and the velocity association to be -0.001 [-0.295, 0.254]. That is even though the schedule of biopsies depended upon observed PSA values it did not lead to a spurious association.

With regards to the second question about the use of right censoring in the simulation study instead of interval censored data, this cannot lead to any differences in results. The reason is that we use a full likelihood approach as described in Section 2 of the original manuscript. Parameter estimation using full likelihood approaches always gives consistent and asymptotically unbiased results (Gentleman and Geyer, 1994).

5. The two sets of AUC's were mistakenly swapped while creating the table and hence the counterintuitive results. We have corrected this mistake, and in addition as advised by the referee, we have also reported the confidence interval. To obtain the 95% confidence interval, we used 15 bootstrapped datasets of patients from the training dataset utilized in this paper. We calculate the AUC at year one, year two and year three of follow-up in AS. The time window for which the AUC is calculated is one year. The resulting estimates are presented in Table 2 below, as well as added in Supplementary material (Web Table 3).

Table 2: Area under the receiver operating characteristic curves, and 95% confidence interval in brackets, obtained using 15 bootstrapped datasets.

Year	log ₂ PSA value and velocity association	log ₂ PSA value association
1	0.613 [0.582, 0.632]	0.595 [0.565, 0.618]
2	0.648 [0.608, 0.685]	0.609 [0.568, 0.654]
3	0.593 [0.560, 0.638]	0.590 [0.536, 0.628]

6. We thank the referee for pointing out these errors. We have fixed these in the updated manuscript.

Response to 2nd Referee's Comments

We would like to thank the Referee for his/her constructive comments, which have allowed us to considerably improve our paper. The main differences of the new version of the manuscript compared to the previous one can be found in Sections 2, 3 and 5. In addition, changes regarding the specific comments have been made throughout the text. You may find below our response to the specific issues raised.

- 1,4. As the referee noted, the equation for the longitudinal sub-model on page 4 of the original manuscript does not indicate that we used a log transform for PSA levels. This is however the general form of the equation for the longitudinal sub-model, and is only used to introduce the joint model (JM) notation. The actual equation, showing the log transformed PSA levels, baseline covariates and B-spline for the effect of time is Equation (7) in the original manuscript. That is, it is not the case that the log transformation is used only in simulation study as noted by the referee, but also used for fitting the PRIAS data.

Since concerns regarding assumption of normality on errors were also raised by the first referee, we refitted our model with an assumption that the errors are T-distributed (df=3). The residual quantile-quantile plots for the model with normally distributed errors as well as the T-distributed (df=3) errors are shown in Figure 1. In addition, the fitted marginal \log_2 PSA profiles, and subject specific fitted as well as observed \log_2 PSA profiles of 9 randomly selected patients, using the two different models are presented in Figure 2 and Figure 3, respectively.

With regards to the fitted profiles for the 3 demonstration patients, we show their fitted profiles in Figure 7. The fitted profiles are dynamic in nature, and utilize information from both the observed PSA levels and time of latest biopsy. The two time points for each of the patients are the same time points at which we made personalized schedules for these patients in the original manuscript.

2. The referee noted that we define $\mathcal{M}_i(t)$ as PSA level up-to last biopsy, and it should not be the case. However on page 5 of the original manuscript it is instead defined as the following: "... where $\mathcal{M}_i(t) = \{m_i(v), 0 \leq v \leq t\}$ denotes the history of the underlying PSA levels up to time t ". This is a standard notation can also be found in existing literature (Rizopoulos, 2012; Tsiatis and Davidian, 2004).

Regarding the second comment about ability to use historical PSA measurements to make

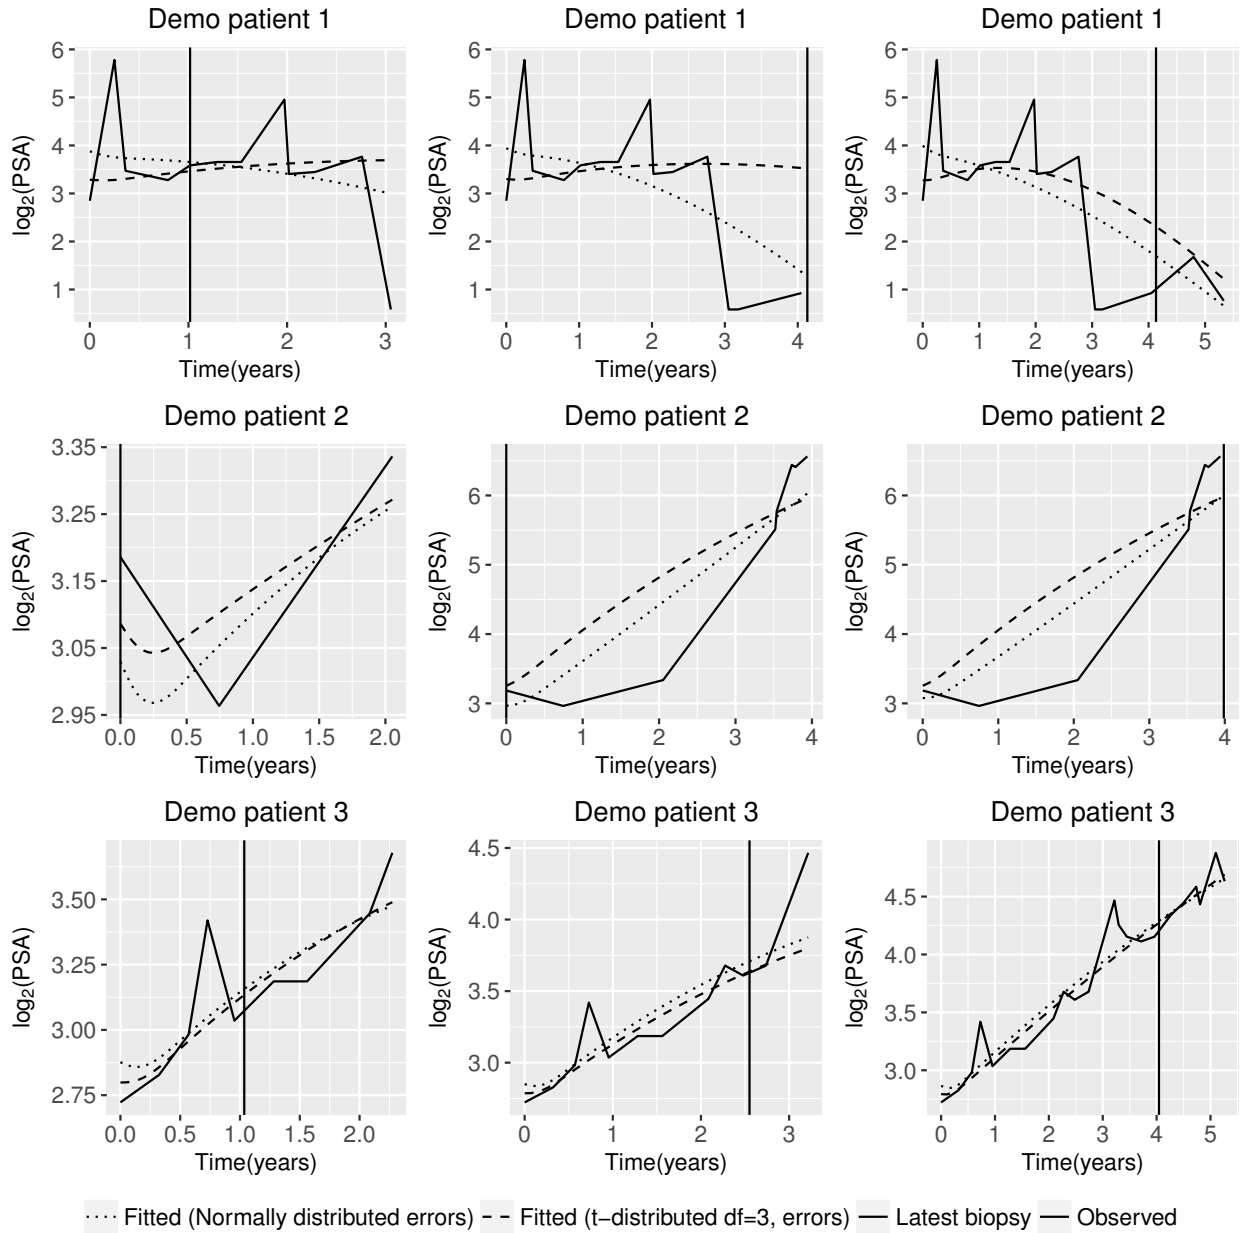


Figure 7: Fitted versus observed \log_2 PSA profiles for the three demonstration patients, at two different time points. The vertical line with a dot-dash pattern shows the time of the latest biopsy. The fitted profiles are dynamic in nature, and utilize information from both the observed PSA levels and time of latest biopsy.

a decision about immediate/delayed biopsy, yes we indeed provide a method to “evaluate biopsy time from current time, particularly when there is new information, such as new PSA measure after last biopsy”. To illustrate this, suppose for the j -th patient, the last biopsy was conducted at time t , and the current visit time at which PSA is measured is $s > t$, then we are interested in finding the time $u > s$ of the next biopsy which utilizes all the available information up to s . To this end, all of our approaches are based on the posterior predictive distribution of GR time, given by $p\{T_j^* \mid T_j^* > t, \mathcal{Y}_j(s), \mathcal{D}_n\}$. Here $\mathcal{Y}_j(s)$ is the history of PSA up to s and the information that no GR was found at last biopsy is included via the condition $T_j^* > t$. Indeed as the referee noted it is possible that $t < T_j^* \leq s$, and then biopsy should be conducted immediately. However, it is often the case that difference between consecutive biopsies is required to be at least an year. Thus even if the schedule also suggests a time $t < u \leq s$, the biopsy should be conducted with a delay of $1 - t$. We have discussed this scenario in Section 3.4 of the original manuscript. In addition, we have shown the entire decision making process related to conducting a biopsy in the flowchart in Figure 1 of the original manuscript.

3. We observed in Figure 2 of the original manuscript that the variance of posterior predictive distribution of event time decreases as more information is gathered over time. That is, a schedule based on expected/median time of Gleason reclassification (GR) is less accurate (the consistency property) in predicting true event time when less information is available. In comparison, the schedule based on dynamic risk of GR is robust in the sense that it is more risk averse than the schedule based on median time of GR (50% risk), at all time points. For example, in PRIAS, on average it schedules biopsies whenever the risk increases more than 5.3%. Thus, it is less likely to overshoot the true GR time by a big margin even if less information is available for the patients. This is also demonstrated via the simulation study, wherein the schedule based on dynamic risk of GR leads to almost the same mean offset and variance of offset across the three subgroups of patients. However, given that the term robust implies a different meaning in a statistical context, we have removed it in the updated manuscript.

Minor Concerns Shared by the 2nd Referee

1. We have updated the captions of tables and graphs in the updated version of the manuscript.
- 2,4. For the two parameters κ and η mentioned by the referee, we do not use fixed set of values.

We compute the parameter κ (dynamic risk of GR) from the data, as shown in Section 3.3 of the original manuscript. That is, we obviate choosing this value manually. When it is chosen manually, it may not always give the best results. For example, often clinicians use a $\kappa = 0.05$ or 5% risk, however as shown in Web Table 4 in supplementary material the performance of this schedule is exactly same as that of annual schedule. That is, it gives a very small offset at the cost of too many biopsies.

With regards to the choice of weights η_1, η_2 , as discussed in Section 4.1 of the original manuscript, this choice can be obviated by reformulating the optimization of original weighted sum as a constrained optimization problem. For example, if η_1 is the weight corresponding to average number of biopsies $E(N^S)$ and η_2 is the weight corresponding to average offset $E(O^S)$, then we can instead put a constraint C on average offset, and then optimize for only the number of biopsies. Since multiple studies have reported small prostate cancer specific mortality in low risk patients enrolled in active surveillance programs (Klotz et al., 2009; Loeb et al., 2016; Tosoian et al., 2011), a recommended cutoff C on average offset is 1 year. We have also added this information in the discussion section of the updated manuscript.

3. We agree to the referee with regards to not having a direct interpretation of age since we also have quadratic form of age in the model. YES WE DID A MISTAKE...DID WE? atleast one of the two is significant anyway

However in the original manuscript we did not interpret the effect size, but rather only mentioned that "... and the age at the time of inclusion in AS were strongly associated with the hazard of GR". This is still true, because both the linear and quadratic effects of age are significant. -WE CALL IT SIGNIFICANT NOW..CHANGE OF SENTENCE

4. Good and bad

References

- Gentleman, Robert and Charles J Geyer (1994). “Maximum likelihood for interval censored data: Consistency and computation”. In: *Biometrika* 81.3, pp. 618–623.
- Huang, Xianzheng, Leonard A Stefanski, and Marie Davidian (2009). “Latent-model robustness in joint models for a primary endpoint and a longitudinal process”. In: *Biometrics* 65.3, pp. 719–727.
- Klotz, Laurence et al. (2009). “Clinical results of long-term follow-up of a large, active surveillance cohort with localized prostate cancer”. In: *Journal of Clinical Oncology* 28.1, pp. 126–131.
- Loeb, Stacy et al. (2016). “Immediate versus delayed prostatectomy: Nationwide population-based study”. In: *Scandinavian journal of urology* 50.4, pp. 246–254.
- Proust-Lima, Cécile et al. (2014). “Joint latent class models for longitudinal and time-to-event data: A review”. In: *Statistical methods in medical research* 23.1, pp. 74–90.
- Razali, Ahmad Mahir and Ali A Al-Wakeel (2013). “Mixture Weibull distributions for fitting failure times data”. In: *Applied Mathematics and Computation* 219.24, pp. 11358–11364.
- Rizopoulos, Dimitris (2012). *Joint Models for Longitudinal and Time-to-Event Data: With Applications in R*. CRC Press.
- Rizopoulos, Dimitris, Geert Verbeke, and Geert Molenberghs (2008). “Shared parameter models under random effects misspecification”. In: *Biometrika* 95.1, pp. 63–74.
- Tosoian, Jeffrey J et al. (2011). “Active surveillance program for prostate cancer: an update of the Johns Hopkins experience”. In: *Journal of Clinical Oncology* 29.16, pp. 2185–2190.
- Tsiatis, Anastasios A and Marie Davidian (2004). “Joint modeling of longitudinal and time-to-event data: an overview”. In: *Statistica Sinica* 14.3, pp. 809–834.